

Diabetes Prediction using Supervised Learning Techniques

By
Oluyomi Alabi

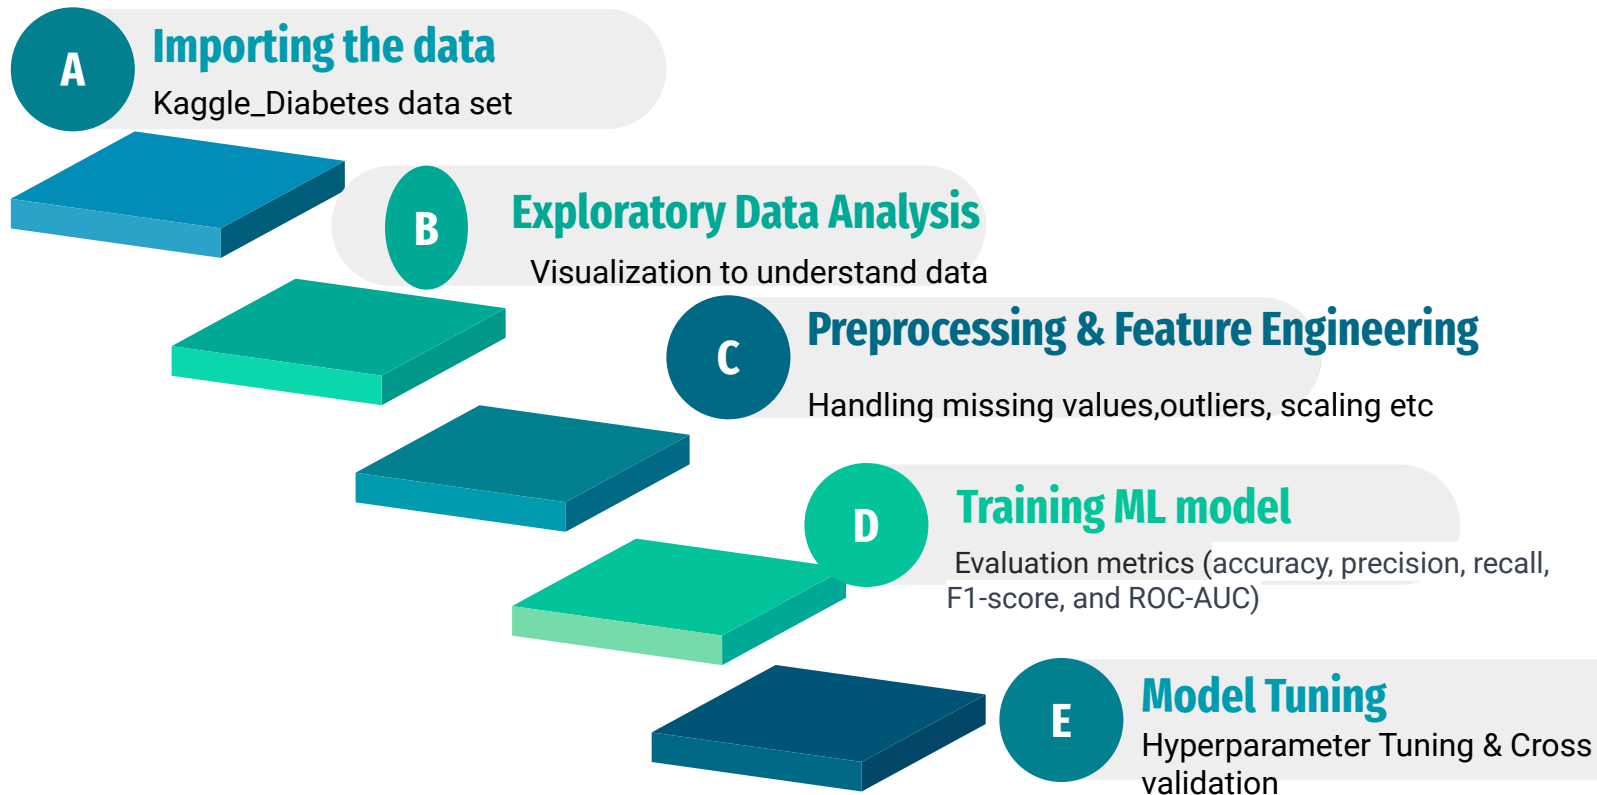
Introduction

This project applied supervised learning techniques to a real-world "Diabetes" dataset from the National Institute of Diabetes and Digestive and Kidney Diseases, and use data visualization tools to communicate the insights gained from the analysis. The objective of the dataset is to diagnostically predict whether a patient has diabetes based on certain diagnostic measurements included in the dataset.

Project Goals

The ultimate goal of the project is to gain insights from the data sets and communicate these insights to stakeholders using appropriate visualizations and metrics to make informed decisions based on the business questions asked.

Process

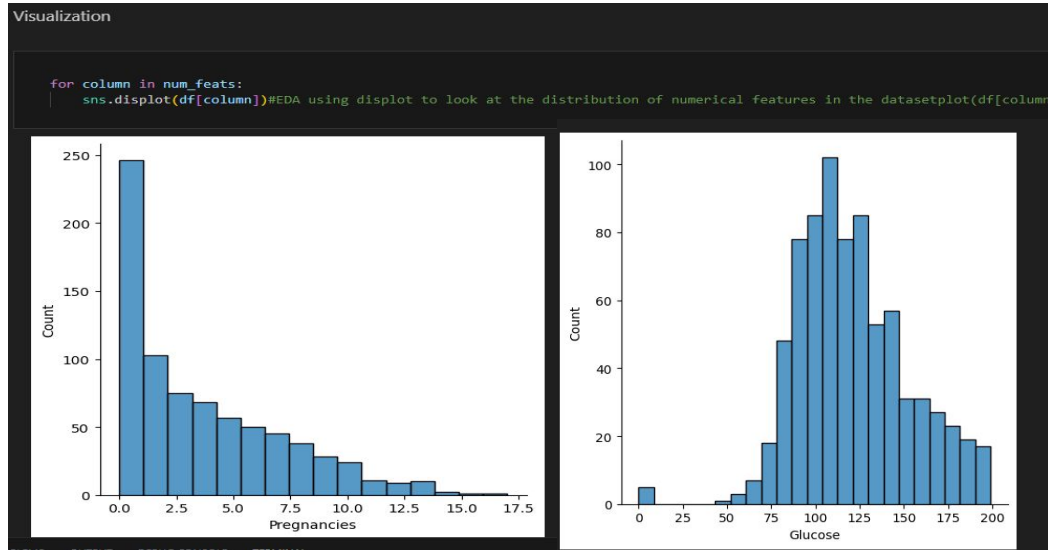


EDA QUESTIONS ANSWERED.

Are there any missing values in the dataset?

- How are the predictor variables related to the outcome variable?
- What is the correlation between the predictor variables?
- What is the distribution of each predictor variable?
- Are there any outliers in the predictor variables?
- How are the predictor variables related to each other?
- Is there any interaction effect between the predictor variables?
- What is the average age of the individuals in the dataset?
- What is the average glucose level for individuals with diabetes and without diabetes?
- What is the average BMI for individuals with diabetes and without diabetes?
- How does the distribution of the predictor variables differ for individuals with diabetes and without diabetes?
- Are there any differences in the predictor variables between males and females (if gender information is available)?

Exploratory Data Analysis Visualization

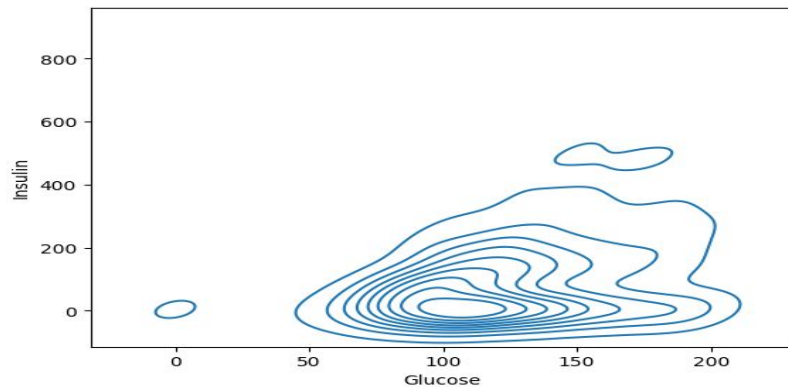


Displot showing distribution of pregnancy and Glucose variables..

Visualizations showing relationship between Glucose level and Insulin.

```
import matplotlib.pyplot as plt
sns.kdeplot(data=df,
            x='Glucose',
            y='Insulin')
plt.suptitle("Kernal Density Plot Comparing Glucose to Insulin")
plt.show()
```

Kernal Density Plot Comparing Glucose to Insulin



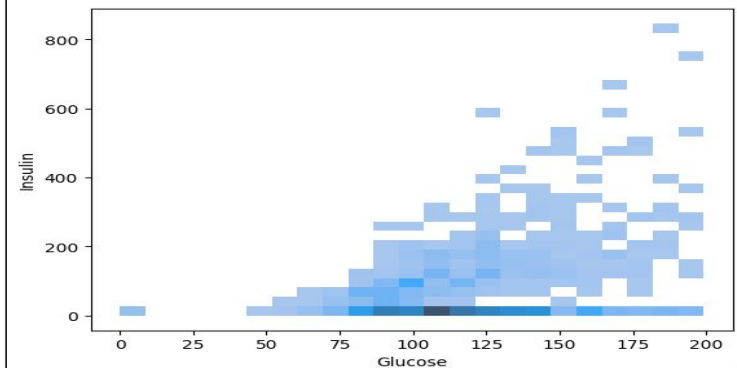
```
import seaborn as sns
import matplotlib.pyplot as plt
sns.histplot(data=df,
             x="Glucose",
             y="Insulin")

plt.suptitle("2 Variable Histogram = Density Plot")

plt.show()
```

✓ 0.3s


2 Variable Histogram = Density Plot



Kernel Density plot and Histogram Density plot comparing Glucose to insulin showing outliers.

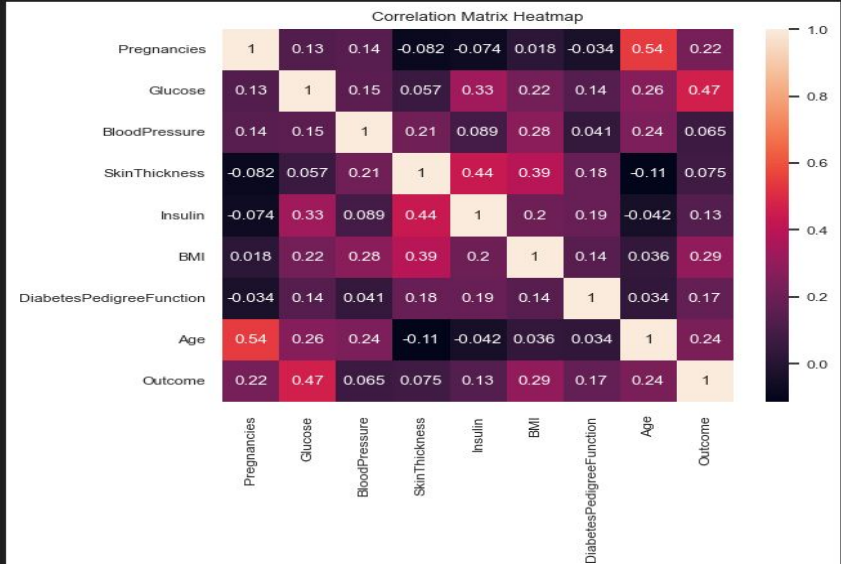
Scattered plot and Correlation matrix map showing correlation between different variables.

```
sns.scatterplot(data=df,  
                x="BloodPressure",  
                y="Age",  
                alpha = 0.7)  
plt.suptitle("Scatter Plot BloodPressure to Age")  
plt.show()
```



```
sns.heatmap(data= corr, annot=True)  
plt.title("Correlation Matrix Heatmap")
```

Text(0.5, 1.0, 'Correlation Matrix Heatmap')



Predictor variables Average grouped by Outcome showing individuals with or without Diabetes.

```
import pandas as pd
# define the columns to be grouped and calculated
columns_to_groupby = ['Outcome']
columns_to_calculate = ['SkinThickness', 'Insulin', 'Glucose', 'BloodPressure', 'BMI', 'DiabetesPedigreeFunction', 'Age']

# Calculate the average (mean) values for the specified columns for each group
Predictorvariables_average_results = df.groupby(columns_to_groupby)[columns_to_calculate].mean().reset_index()

print(Predictorvariables_average_results)
```

| | Outcome | SkinThickness | Insulin | Glucose | BloodPressure | BMI \ |
|---|---------|---------------|------------|------------|---------------|-----------|
| 0 | 0 | 19.664000 | 68.792000 | 109.980000 | 68.184000 | 30.304200 |
| 1 | 1 | 22.164179 | 100.335821 | 141.257463 | 70.824627 | 35.142537 |

| | DiabetesPedigreeFunction | Age |
|---|--------------------------|-----------|
| 0 | 0.429734 | 31.190000 |
| 1 | 0.550500 | 37.067164 |

From this column calculation grouping by Outcome, Average glucose level of individuals without diabetes is 109.98 while the average glucose level with diabetes is 141.257463. Average BMI for individuals without Diabetes is 30.30 while Average BMI of individuals with Diabetes is 35.14

Distribution statistics showing statistical significant correlation between the feature variables

```
#Distribution statistics showing if there is a statistical significant correlation  
# Calculate the correlation matrix  
corr = df.corr()  
corr
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|--------------------------|-------------|----------|---------------|---------------|-----------|----------|--------------------------|-----------|----------|
| Pregnancies | 1.000000 | 0.129459 | 0.141282 | -0.081672 | -0.073535 | 0.017683 | -0.033523 | 0.544341 | 0.221898 |
| Glucose | 0.129459 | 1.000000 | 0.152590 | 0.057328 | 0.331357 | 0.221071 | 0.137337 | 0.263514 | 0.466581 |
| BloodPressure | 0.141282 | 0.152590 | 1.000000 | 0.207371 | 0.088933 | 0.281805 | 0.041265 | 0.239528 | 0.065068 |
| SkinThickness | -0.081672 | 0.057328 | 0.207371 | 1.000000 | 0.436783 | 0.392573 | 0.183928 | -0.113970 | 0.074752 |
| Insulin | -0.073535 | 0.331357 | 0.088933 | 0.436783 | 1.000000 | 0.197859 | 0.185071 | -0.042163 | 0.130548 |
| BMI | 0.017683 | 0.221071 | 0.281805 | 0.392573 | 0.197859 | 1.000000 | 0.140647 | 0.036242 | 0.292695 |
| DiabetesPedigreeFunction | -0.033523 | 0.137337 | 0.041265 | 0.183928 | 0.185071 | 0.140647 | 1.000000 | 0.033561 | 0.173844 |
| Age | 0.544341 | 0.263514 | 0.239528 | -0.113970 | -0.042163 | 0.036242 | 0.033561 | 1.000000 | 0.238356 |
| Outcome | 0.221898 | 0.466581 | 0.065068 | 0.074752 | 0.130548 | 0.292695 | 0.173844 | 0.238356 | 1.000000 |

Summary statistics of each variable showing anomalies and potential outliers

```
#The summary statistics of each of the variables, we can identify anomalies and potential outliers.  
df.describe()
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|-------------|------------|---------------|---------------|------------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

E.g BloodPressure cannot be zero except the patient is dead, same with Skin Thickness, BMI where minimum value is zero

This summary statistics shows the average age of the individuals in the dataset as 33 years, Overall Average glucose level as 120.89, Average BMI as 31.99.

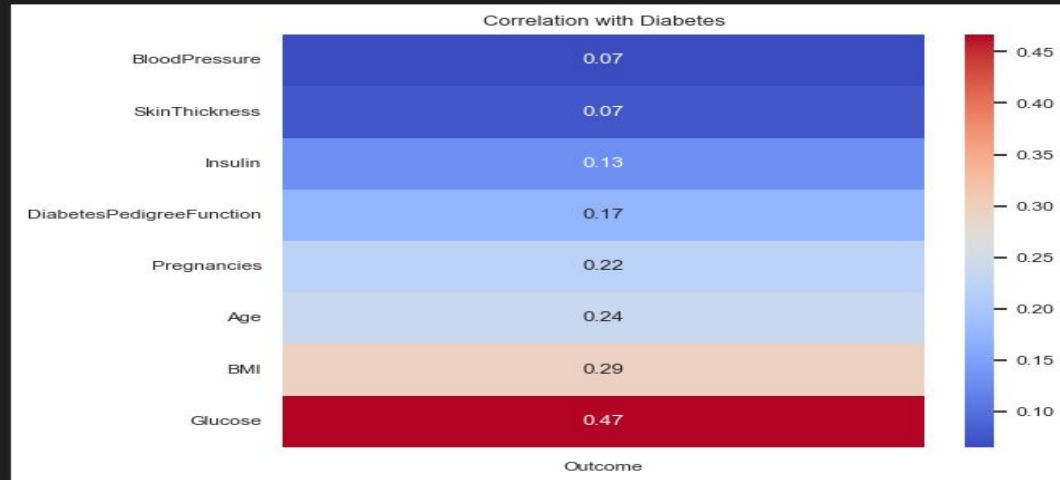
This statistics also show some anomalies eg Blood Pressure of a living being cannot be zero, same with Skin Thickness, BMI, Glucose level. This indicates outliers.

Various feature variable correlation indicating Diabetes in target variable(Outcome).

```
corr = df.corr()
target_corr = corr['Outcome'].drop('Outcome')

target_corr_sorted = target_corr.sort_values(ascending=True)

sns.set(font_scale=0.8)
sns.set_style("white")
sns.set_palette("PuBuGn_d")
sns.heatmap(target_corr_sorted.to_frame(), cmap="coolwarm", annot=True, fmt='.2f')
plt.title('Correlation with Diabetes')
plt.show()
```



With a correlation of 0.47, Glucose is the most strongly correlated feature with the outcome. This suggests that higher glucose levels could be a significant indicator of diabetes.

BMI can also be an important factor in having diabetes.

Results showing no missing values in the dataset.

```
# check missing values in variables  
df.isnull().sum()  
✓ 0.0s
```

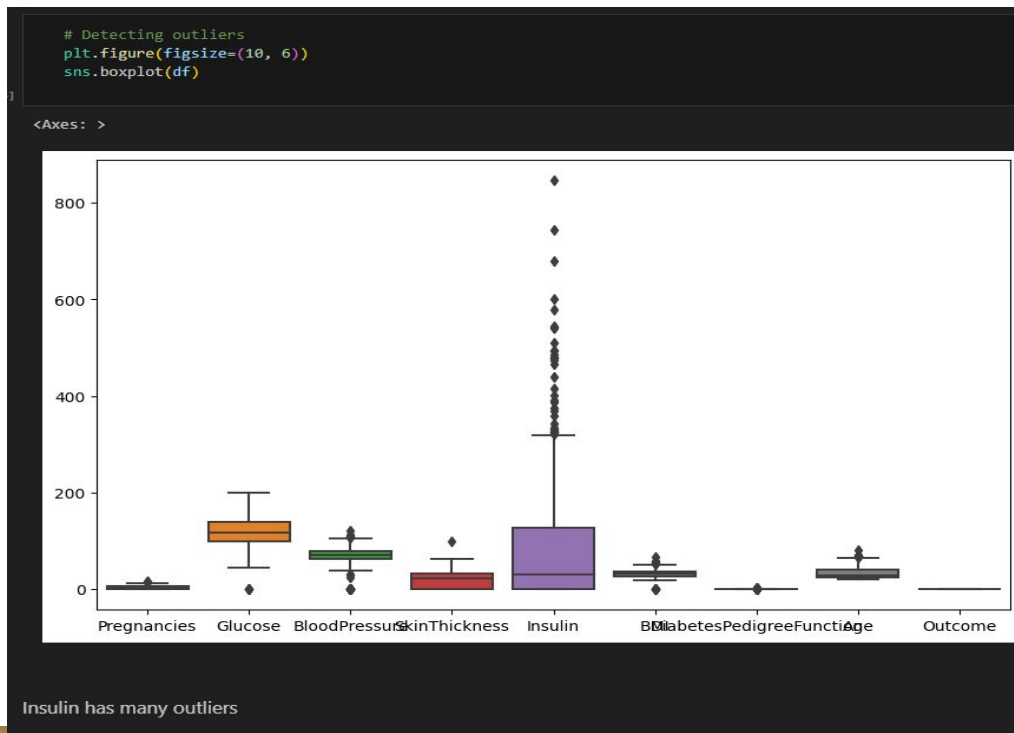
| | |
|--------------------------|---|
| Pregnancies | 0 |
| Glucose | 0 |
| BloodPressure | 0 |
| SkinThickness | 0 |
| Insulin | 0 |
| BMI | 0 |
| DiabetesPedigreeFunction | 0 |
| Age | 0 |
| Outcome | 0 |

dtype: int64

Answer: There are no missing values in the dataset.

There are no missing values in the dataset

Handling outliers in preprocessing and feature engineering



It can be seen that Insulin has many outliers more than other feature variables, this was handled by transforming using log transformation.

```
cols_with_outliers = ['Glucose', 'Insulin',  
                       'SkinThickness', 'BloodPressure', 'BMI',  
                       'Age']  
df[cols_with_outliers] =  
df[cols_with_outliers].apply(lambda x:  
                               np.log1p(x))
```

Logistic regression model showing evaluation metrics results

```
# Evaluate the model using various metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
confusion_mat = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Evaluation Metrics:")
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
print(f"ROC-AUC: {roc_auc}")
print(f"Confusion Matrix:\n{confusion_mat}")
print(f"classification_report: {classification_rep}")
```

✓ 0.0s

```
Evaluation Metrics:
Accuracy: 0.7272727272727273
Precision: 0.6164383561643836
Recall: 0.5625
F1 Score: 0.5882352941176471
ROC-AUC: 0.6885347682119204
Confusion Matrix:
[[123  28]
 [ 35  45]]
```

Evaluation Metrics:

Accuracy: 0.7272727272727273

Precision: 0.6164383561643836

Recall: 0.5625

F1 Score: 0.5882352941176471

ROC-AUC: 0.6885347682119204

Confusion Matrix:

[[123 28]

[35 45]]

| classification_report: | | precision | recall | f1-score | support |
|------------------------|--|-----------|--------|----------|---------|
|------------------------|--|-----------|--------|----------|---------|

| | | | | |
|---|------|------|------|-----|
| 0 | 0.78 | 0.81 | 0.80 | 151 |
|---|------|------|------|-----|

| | | | | |
|---|------|------|------|----|
| 1 | 0.62 | 0.56 | 0.59 | 80 |
|---|------|------|------|----|

| | | | |
|----------|--|------|-----|
| accuracy | | 0.73 | 231 |
|----------|--|------|-----|

| | | | | |
|-----------|------|------|------|-----|
| macro avg | 0.70 | 0.69 | 0.69 | 231 |
|-----------|------|------|------|-----|

| | | | | |
|--------------|------|------|------|-----|
| weighted avg | 0.72 | 0.73 | 0.72 | 231 |
|--------------|------|------|------|-----|

Random Forest Classifier model with model Evaluation metrics Results

```
# RandomForestClassifier model evaluation metrics
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
```

```
print('Model accuracy score with 10 decision-trees : {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
print("ROC-AUC:", roc_auc)
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_rep)
```

✓ 0.1s

Model accuracy score with 10 decision-trees : 0.7316
Precision: 0.6097560975609756
Recall: 0.625
F1-score: 0.6172839506172839
ROC-AUC: 0.7065397350993378

Confusion Matrix:
[[119 32]
 [30 50]]

Model accuracy score with 10 decision-trees : 0.7316

Precision: 0.6097560975609756

Recall: 0.625

F1-score: 0.6172839506172839

ROC-AUC: 0.7065397350993378

Confusion Matrix:

```
[[119 32]
 [ 30 50]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.80 | 0.79 | 0.79 | 151 |
| 1 | 0.61 | 0.62 | 0.62 | 80 |
| accuracy | | | 0.73 | 231 |
| macro avg | 0.70 | 0.71 | 0.71 | 231 |
| weighted avg | 0.73 | 0.73 | 0.73 | 231 |

Result explanation

Based on these metrics, we can Random Forest Classifier model has a slightly higher accuracy score (0.7316), compared to Logistic Regression model with .07272 accuracy. The Random Forest Classifier has a higher ROC-AUC score (0.7065) compared to the Logistic Regression model (0.6885). Additionally, the Random Forest Classifier has slightly higher F1 score and recall values.

Overall, the Random Forest Classifier appears to perform slightly better than the Logistic Regression model based on these evaluation metrics.

model tuning and cross-validation, to improve and optimize the model's performance.

```
#Model performance evaluation matrices
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

# Print the evaluation metrics
print("Best Hyperparameters:", best_params)
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
print("ROC-AUC:", roc_auc)
```

✓ 0.0s

```
Best Hyperparameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 50}
Accuracy: 0.7402597402597403
Precision: 0.6190476190476191
Recall: 0.65
F1-score: 0.6341463414634146
ROC-AUC: 0.7190397350993377
```

This was done to improve and optimize the RandomForestClassifier model's performance. By Hyperparameter Tuning the number of estimators to 50. Evaluation metrics accuracy Increased.

Accuracy: 0.74025974
Precision: 0.61904761
Recall: 0.65
F1-score: 0.63414634
ROC-AUC: 0.71903973
Mean Cross-Validation Accuracy: 0.7820353
Mean cv precision: 0.72295

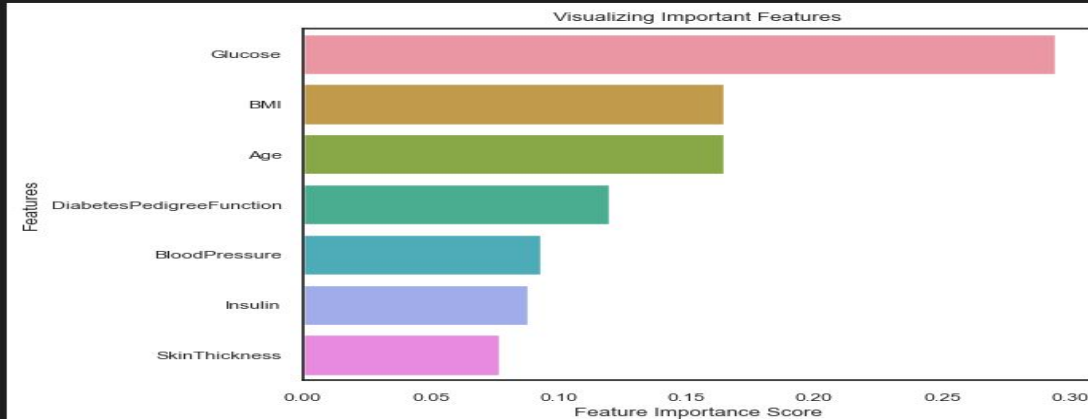
Feature scores showing the order of Importance or the diagnostic variables after tuning

```
sns.barplot(x=feature_scores, y=feature_scores.index)

# Add labels to the graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')

# Add title to the graph
plt.title("Visualizing Important Features")

# Visualize the graph
plt.show()
```



It can be seen that Glucose level is the most important feature in this diagnostic model while skin Thickness is the least feature affecting the accuracy of this model.

conclusion

From the machine learning developed, and exploratory data analysis conducted, the following were observed.

- Random Forest Classifier appear to have a better accuracy than Logistic regression model in predicting whether a patient have Diabetes or not.
- Cross validation gave more robust estimate of the mode performances with increased accuracy as 0.782035306 and precision of 0.72297510.
- Glucose level is the most important feature in predicting whether a patient have Diabetes or not.
- The least important feature that doesn't really affect the prediction is Skin thickness.

Challenges

- Time constraints to explore available data using other models

THANK YOU